

Unveiling the Hidden Echoes: Zero-Dependency Forensic Reconstruction of 'Deleted' LLM Chats from Browser & Desktop LevelDB Instances

Dr. Sapna Vikram Mewundi
Lead Forensics Researcher, Verity Project
sapna@local.dev

ABSTRACT

As Large Language Model (LLM) portals become standard corporate utilities, proprietary source code, credentials, and sensitive configurations are routinely processed by users. When a user clicks 'Delete Chat' inside ChatGPT, Claude, or Gemini interfaces, they expect their local trace data to be permanently erased. However, because Chromium-based browsers (Chrome, Edge) and Electron desktop clients store this telemetry inside IndexedDB databases backed by Google's LevelDB engine, these deleted records persist on disk in write-ahead logs (.log) and uncompactd Sorted String Tables (.sst.ldb). This paper details the reverse engineering of client-side LLM storage schemas and the V8 deserialization format. We present the mechanics of rebuilding conversation trees directly from raw binary fragments. Finally, we introduce a zero-dependency, pure-Python carving methodology to bypass active database locks and dynamically reconstruct deleted chat histories across all user profiles, providing incident responders with a secure, cryptographically validated chain of custody.

I. INTRODUCTION

The security industry has focused heavily on Generative AI security at the boundary layer: prompt injections, data guardrails, and cloud storage compliance. Virtually unmapped, however, is the client-side forensic footprint left behind by web interfaces and native wrapper applications on employee workstations.

When a user deletes a chat thread, the cloud-side storage is updated. Locally on the endpoint, however, the deletion is simply written as a LevelDB tombstone or index update. Because LevelDB is an append-only log-structured merge-tree (LSM tree), the actual serialized data blocks containing the prompts and responses remain intact in write-ahead log records or orphaned data blocks until compaction occurs. This creates a critical forensic window: an attacker or examiner with local access can recover sensitive, supposedly 'deleted' data.

II. LOW-LEVEL DATABASE ARCHITECTURE: INDEXEDDB & LEVELDB

Chromium browsers implement IndexedDB using LevelDB as the underlying storage engine. On Windows, these files reside in user AppData folders:

```
%LOCALAPPDATA%\Google\Chrome\User  
Data\<ProfileName>\IndexedDB\https_chatgpt.com_0.indexeddb.leveldb\
```

Traditional forensics tools rely on loading LevelDB using compiled C++ drivers (like plyvel). During a live incident response triage, installing compiler tools on production machines is prohibited. Incident responders require a zero-dependency byte carver that directly dissects raw files.

III. REVERSING V8 SERIALIZATION: THE BINARY STRUCTURE

Chrome and Edge store IndexedDB records serialized in V8's internal binary serialization format (ValueSerializer).

A. Varint Decoding

V8 uses Google Protocol Buffer style varints (variable-length integers) to encode lengths and indexes. Every byte's highest-order bit (bit 7) is a flag: 1 means another byte follows, and 0 indicates the end of the varint. The lower 7 bits of each byte are concatenated to form the integer value.

B. V8 String Tags and UTF-16LE Encoding

Strings in the serialized format are prefixed with data type tags: OneByteString (0x22) indicates ASCII string, and TwoByteString (0x63) indicates UTF-16LE string. Crucially, the varint length of TwoByteString represents the number of bytes, not characters. The string is decoded by reading exactly length bytes and translating them using UTF-16LE.

IV. MESSAGE REASSEMBLY & ROLE RECONSTRUCTION

To map raw strings into chronological conversations, we reverse-engineered the serialization structure of the messages array.

A. Smi-Shifted Index Keys

V8 serializes array element properties using small integers (Smi) as keys. V8 encodes Smis by shifting the integer left by 1 bit (encoded = actual $\ll$ 1). The actual message index indicates the role: Odd actual indices represent User Prompts, and Even actual indices represent Assistant Responses. The role mapping is mathematically defined as:

```
role = 'user' if (index / 2) mod 2 = 1 else 'assistant'
```

B. Nesting Depth Tracking

V8 serialized message objects contain nested objects. Because object ends are marked by 0x7b, a simple parser would break out of the message parsing loop prematurely. A forensic-grade parser must track nesting depth, incrementing on object (0x6f) or array (0x61) tags, and decrementing on object end (0x7b).

V. THE VERITY FORENSIC FRAMEWORK

To operationalize these findings, we developed Verity, a zero-dependency, pure-Python carving tool. It bypasses Windows exclusive write locks by accessing write-ahead logs in read-only binary mode, cloning files dynamically, and reassembling conversation threads across all user profiles. To ensure evidence integrity, Verity seals the output using HMAC-SHA256 signatures, validated in-browser via the Web Cryptography API.

VI. CONCLUSION

This paper demonstrates that client-side ephemeral AI sessions are a myth. By parsing raw V8 serialized objects from LevelDB logs, investigators can fully recover deleted records. The Verity framework provides incident responders with a secure, lightweight tool to audit endpoints and protect corporate data assets.